




# **100 SPORTS IMAGE CLASSIFICATION USING EFFICIENTNETB0 AND RESNET50**

**Exploring Ensemble  
Learning and Fine-  
Tuning**  
**By**  
**John Yare**  
**& Joshua Adegbola**

# OVERVIEW

- ❖ Objective: Classify 100 sports image categories using an ensemble model by combining predictions of a Vision Transformer alternative model available in keras (EfficientNetB0) and ResNet50 model.
- ❖ Models: EfficientNetB0 and ResNet50.
- ❖ Dataset:
  - Initial manipulated dataset had less training data.
  - Reverted to baseline dataset for better performance and more variety of data in training.
- ❖ Achievements:
  - Achieved 92.6% test accuracy with EfficientNetB0 model.
  - Achieved 96.2% test accuracy with ResNet50 model.
  - Combined predictions from EfficientNetB0 and ResNet50 by optimizing weights for each model, based on each model's validation performances, into an ensemble model to improve the model's performance.
  - Achieved 90% test accuracy with ensemble model.

# BLOCK DIAGRAMS



Block Diagram: Dataset Loading to Ensemble Evaluation

# DATASET PREPARATION

- ❖ Dataset paths for training, validation, and testing.
- ❖ Used Label Encoder for class label encoding.
- ❖ Visualized class distribution in training data.



Bar Chart: Image Count Per Class

# DATA PREPROCESSING AND AUGMENTATION

```
# Data augmentation layer for applying random transformations to input images
augment = tf.keras.Sequential([
    RandomRotation(0.1), # Randomly rotate images by 10%
    RandomZoom(0.1), # Randomly zoom into images
    RandomContrast(0.1), # Randomly adjust image contrast
    RandomFlip("horizontal_and_vertical") # Randomly flip images horizontally and vertically
], name='AugmentationLayer')

# Define the input layer for the model
inputs = layers.Input(shape = (224,224,3), name='inputLayer')

# Pass inputs through the augmentation layer
x = augment(inputs)

# Pass the augmented inputs through the pretrained model
pretrain_out = pretrained_model(x, training = False)

# Add a fully connected dense layer for feature extraction
x = layers.Dense(350)(pretrain_out)
x = layers.Activation(activation="relu")(x) # Apply ReLU activation
x = BatchNormalization()(x) # Normalize activations to speed up training
x = layers.Dropout(0.25)(x) # Add dropout to reduce overfitting

# Add an output layer with the number of classes (num_classes)
x = layers.Dense(num_classes)(x)
outputs = layers.Activation(activation="softmax", dtype=tf.float32, name='activationLayer')(x) # mixed_precision need separated Dense and Activation layers

# Create the final model
model = Model(inputs=inputs, outputs=outputs)
```

Diagram: Augmentation Applied to Training Data Only

❖ Challenges with direct augmentation:

- Reduced training and validation accuracy.
- Increased losses indicated no learning.

❖ Solution:

- Augmentation applied after the input layer for training data only.

# EFFICIENTNETB0 MODEL ARCHITECTURE

- ❖ Base Model: Pretrained on ImageNet.
- ❖ Custom Layers:
  - Dense layers for feature extraction.
  - Dropout for regularization.
  - Batch Normalization for stability.
  - Softmax activation for classification.
- ❖ Fine-tuning with frozen Batch Normalization layers to avoid instability.

Layer (type)	Output Shape	Param #
inputLayer (InputLayer)	(None, 224, 224, 3)	0
cast_1 (cast)	(None, 224, 224, 3)	0
AugmentationLayer (Sequential)	(None, 224, 224, 3)	0
efficientnetb0 (Functional)	(None, 1280)	4,049,571
dense (Dense)	(None, 350)	448,350
activation (Activation)	(None, 350)	0
batch_normalization (BatchNormalization)	(None, 350)	1,400
dropout (Dropout)	(None, 350)	0
dense_1 (Dense)	(None, 100)	35,100
cast_2 (cast)	(None, 100)	0
activationLayer (Activation)	(None, 100)	0
Total params:	4,534,421	(17.30 MB)
Trainable params:	484,150	(1.85 MB)
Non-trainable params:	4,050,271	(15.45 MB)

Diagram: Model summary

# TRAINING THE EFFICIENTNETB0 MODEL

## ◆ Training settings:


- Optimizer: Adam (learning rate: 0.0005).
- Batch size: 10.
- Epochs: 10.

## ◆ Callbacks:

- EarlyStopping to prevent overfitting.
- ReduceLROnPlateau for dynamic learning rate adjustment.

## ◆ Results:


- Improved training and validation performance.



Plots: Training and Validation Loss/Accuracy

# FINE-TUNING THE EFFICIENTNETB0 MODEL

- ❖ Unfroze EfficientNetB0 layers while freezing BatchNormalization.
- ❖ Reduced learning rate to 0.0001 for stability.
- ❖ Reused EarlyStopping and ReduceLROnPlateau callbacks.
- ❖ Results:
  - Further improvement in validation accuracy and loss reduction.



## Test accuracy performance for EfficientNetB0 model

### Evaluating the Model

Evaluate the model on the test set

```
[ ] results = model.evaluate(test_images, verbose=0)  
print("    Test Loss: {:.5f}".format(results[0]))  
print("Test Accuracy: {:.2f}%".format(results[1] * 100))
```

```
→ Test Loss: 0.28631  
Test Accuracy: 92.60%
```

# MODEL ARCHITECTURE: RESNET50

Base Model:  
ResNet-50  
pretrained on  
ImageNet.

Output Layer:  
Dense Layer with  
100 softmax nodes  
for classification

Freezing ResNet-50  
layers for transfer  
learning.

```
Building, Compiling and Training the model

[] model1 = Sequential()

model1.add(ResNet50(include_top = False, pooling = 'avg',
                    weights = resnet_weights_path))

[] Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94765736/94765736 5s @us/step

[] model1.add(Dense(100, activation = 'softmax'))

[] model1.summary()
Model: "sequential"
+---+
| Layer (type)        | Output Shape   | Param # |
+---+
| resnet50 (Functional)| (None, 2048)    | 23,587,712|
| dense_2 (Dense)     | (None, 100)     | 204,000   |
+---+

Total params: 23,792,612 (90.76 MB)
Trainable params: 23,739,492 (90.56 MB)
Non-trainable params: 53,120 (207.50 KB)
```

# TRAINING SETUP

## Loss

**Function:** Categorical Crossentropy.

**Optimizer:** SGD  
(Stochastic Gradient Descent).

**Metrics:** Accuracy.


Training for 10 epochs using train and validation datasets.

```
[ ] model1.compile(optimizer= 'sgd',
                  loss = 'categorical_crossentropy',
                  metrics = ['accuracy'])

[ ] history1 = model1.fit(train_generator,
                         validation_data = val_generator, epochs = 10)

[ ] Epoch 1/10
[ ] 675/675 - 75s 91ms/step - accuracy: 0.4013 - loss: 2.7818 - val_accuracy: 0.8600 - val_loss: 0.5968
[ ] Epoch 2/10
[ ] 675/675 - 52s 76ms/step - accuracy: 0.8665 - loss: 0.5784 - val_accuracy: 0.9020 - val_loss: 0.3903
[ ] Epoch 3/10
[ ] 675/675 - 52s 76ms/step - accuracy: 0.9206 - loss: 0.3680 - val_accuracy: 0.9260 - val_loss: 0.3090
[ ] Epoch 4/10
[ ] 675/675 - 51s 75ms/step - accuracy: 0.9440 - loss: 0.2819 - val_accuracy: 0.9120 - val_loss: 0.2847
[ ] Epoch 5/10
[ ] 675/675 - 52s 76ms/step - accuracy: 0.9602 - loss: 0.2244 - val_accuracy: 0.9220 - val_loss: 0.2616
[ ] Epoch 6/10
[ ] 675/675 - 53s 78ms/step - accuracy: 0.9705 - loss: 0.1863 - val_accuracy: 0.9360 - val_loss: 0.2423
[ ] Epoch 7/10
[ ] 675/675 - 52s 76ms/step - accuracy: 0.9738 - loss: 0.1628 - val_accuracy: 0.9340 - val_loss: 0.2186
[ ] Epoch 8/10
[ ] 675/675 - 52s 76ms/step - accuracy: 0.9832 - loss: 0.1368 - val_accuracy: 0.9340 - val_loss: 0.2118
[ ] Epoch 9/10
[ ] 675/675 - 52s 76ms/step - accuracy: 0.9868 - loss: 0.1197 - val_accuracy: 0.9480 - val_loss: 0.1926
[ ] Epoch 10/10
[ ] 675/675 - 52s 76ms/step - accuracy: 0.9893 - loss: 0.1092 - val_accuracy: 0.9500 - val_loss: 0.1917
```

# RESULT VISUALIZATION



- ❖ Results:
  - Improved training and validation performance accuracy

# MODEL EVALUATION


## Metrics Used

**1. Loss:** Measures the error during the evaluation process.

1. Lower loss values indicate better model predictions.
2. Example from the evaluation:
  1. **Loss = 0.1547**

**2. Accuracy:** Percentage of correct predictions made by the model.

1. High accuracy indicates the model is correctly classifying most of the input images.
2. Example from the evaluation:
  1. **Accuracy = 96.20%**



```
↳ Evaluating the model
[ ] tester = model1.evaluate(test_generator)
print("Loss of the model1 is - " , tester[0])
print("Accuracy of the model1 is - " , tester[1]*100 , "%")
25/25 2s 75ms/step - accuracy: 0.9545 - loss: 0.1657
Loss of the model1 is - 0.15479682385921478
Accuracy of the model1 is - 96.20000123977661 %
```

# COMBINING EFFICIENTNETB0 AND RESNET50 PREDICTIONS

## ❖ Ensemble Method:

- Weighted averaging of predictions from both models.
- EfficientNetB0 and ResNet50 assigned weights based on each model's validation and prediction performances.

## ❖ Results:

- Achieved a high performance from combined predictions in ensemble model.

## ❖ Advantages:

- Leverages strengths of both models.
- Reduces individual model biases and overfitting in general.
- Improves generalization, which makes the combined model likely perform better on unseen data compared to a single model.

# DIAGRAM: ENSEMBLE PREDICTION PROCESS

```
# Assuming `test_images` is already defined and properly preprocessed
resnet_predictions = model_resnet.predict(test_images)
#resnet_predictions = model_resnet.predict(test_images)
efficientnet_predictions = model_efficientnet.predict(test_images)

# Weighted averaging for ensemble predictions
weights_resnet = 0.6
weights_efficientnet = 0.4
ensemble_predictions = (weights_resnet * resnet_predictions) + (weights_efficientnet * efficientnet_predictions)
```

```
from sklearn.metrics import accuracy_score

# Convert averaged predictions into class labels
ensemble_labels = np.argmax(ensemble_predictions, axis=1)

# Ground truth labels from test data
true_labels = test_images.classes

# Calculate ensemble accuracy
ensemble_accuracy = accuracy_score(true_labels, ensemble_labels)
print(f"Ensemble Model Accuracy: {ensemble_accuracy * 100:.2f}%")

# Calculate ensemble f1 score
ensemble_f1_score = f1_score(true_labels, ensemble_labels, average='macro')
print(f"Ensemble Model F1 Score: {ensemble_f1_score:.4f}")

→ Ensemble Model Accuracy: 90.00%
Ensemble Model F1 Score: 0.8976
```

# EVALUATION METRICS FOR ENSEMBLE MODEL


- ❖ Achieved Accuracy: 90%.
- ❖ Metrics explained:
  - Precision: True positives among predicted positives.
  - Recall: True positives among actual positives.
  - F1-Score: Balances precision and recall.
  - Confusion Matrix: Breakdown of true vs. predicted classes.

	precision	recall	f1-score	support
air hockey	1.00	1.00	1.00	5
amputee football	1.00	1.00	1.00	5
archery	0.83	1.00	0.91	5
arm wrestling	1.00	1.00	1.00	5
axe throwing	1.00	1.00	1.00	5
balance beam	1.00	1.00	1.00	5
barell racing	1.00	1.00	1.00	5
baseball	0.88	0.88	0.88	5
basketball	1.00	0.40	0.57	5
baton twirling	1.00	1.00	1.00	5
bike polo	1.00	0.48	0.57	5
billiards	1.00	1.00	1.00	5
bmx	0.67	0.88	0.73	5
bobsled	1.00	0.88	0.89	5
bowling	1.00	0.88	0.89	5
boxing	1.00	0.88	0.89	5
bull riding	1.00	1.00	1.00	5
bungee jumping	1.00	1.00	1.00	5
canoe slamon	1.00	1.00	1.00	5
cheerleading	1.00	0.88	0.89	5
surfing	1.00	0.88	0.89	5
swimming	1.00	1.00	1.00	5
table tennis	1.00	0.60	0.75	5
tennis	1.00	1.00	1.00	5
track bicycle	1.00	1.00	1.00	5
trapeze	0.80	0.80	0.80	5
tug of war	1.00	1.00	1.00	5
ultimate	1.00	0.60	0.75	5
uneven bars	0.83	1.00	0.91	5
volleyball	1.00	1.00	1.00	5
water cycling	1.00	1.00	1.00	5
water polo	0.83	1.00	0.91	5
weightlifting	1.00	1.00	1.00	5
wheelchair basketball	1.00	0.80	0.89	5
wheelchair racing	0.71	1.00	0.83	5
wingsuit flying	1.00	1.00	1.00	5
accuracy			0.90	500
macro avg	0.93	0.90	0.90	500
weighted avg	0.93	0.90	0.90	500

Snippet of Ensemble model Classification Report

# VISUALIZING PREDICTIONS

- ❖ Visualized predictions from EfficientNetB0, ResNet50, and ensemble model.
- ❖ Highlighted misclassifications with high confidence.



Grid: True Labels vs. Predictions

# SUMMARY AND KEY TAKEAWAYS

- ❖ EfficientNetB0 and ResNet50 proved effective for sports image classification.
- ❖ Ensemble learning improved generalization, and robustness.
- ❖ Future Work:
  - Experiment with different ensemble strategies.
  - Test on larger datasets.

# REFERENCES

*100 Sports Image Classification.* (2023, May 3).

Kaggle. <https://www.kaggle.com/datasets/gpiosenka/sports-classification/data>

Affonso, C., Rossi, A. L. D., Vieira, F. H. A., & de Leon Ferreira, A. C. P. (2017). Deep learning for biological image classification. *Expert systems with applications*, 85, 114-122.

Hamedghorbani. (2023b, July 17). *100 Sports Classification - EfficientNetB0 | 97.8%*.

Kaggle. <https://www.kaggle.com/code/hamedghorbani/100-sports-classification-efficientnetb0-97-8>

Ishaparanje. (2023, April 19). *sports classification RESNET transfer learning*.

Kaggle. <https://www.kaggle.com/code/ishaparanje/sports-classification-resnet-transfer-learning/notebook>

Liu, X. (2024). Comparison of Four Convolutional Neural Network-Based Algorithms for Sports Image Classification. In *Advances in intelligent systems research/Advances in Intelligent Systems Research* (pp. 178–186). [https://doi.org/10.2991/978-94-6463-370-2\\_20](https://doi.org/10.2991/978-94-6463-370-2_20)

Podgorelec, V., Pečnik, Š., & Vrbančič, G. (2020). Classification of Similar Sports Images Using Convolutional Neural Network with Hyper-Parameter Optimization. *Applied Sciences*, 10(23), 8494. <https://doi.org/10.3390/app10238494>